

Research Associate KTP with the University of Essex and Croud
Inc
Technical Task

Ashkan Lotfipoor
School of Energy, Geoscience, Infrastructure and Society (EGIS),
Heriot-Watt University

July 26, 2023

Introduction

This report is the result of my analysis of the technical assignment for the KTP Associate with the University of Essex and Croud Inc. This report is divided into two sections. The first section presents the results of the marketing data modelling on the provided dataset. Then, section two provide results for the question in Bayesian methods section of the task.

1 Classical Marketing Data Modelling

1.1 Modelling

In this section several regression models are built to determine the effectiveness of marketing strategies of the company and also to predict the revenue. Firstly the data is divided to two dataset, the train and test set, to effectively validate the results of the model on unseen data. The 2016/05/30 date is used as the split point in the dataset.

In this section, the temporal aspect of the data is not considered in the modelling process. So, the week number and Date column are dropped from the dataframe. The remaining variables or predictors in the data are media1_S (spend on media 1), media2_S (spend on media 2), media3_S (spend on media 3), competitor_sales (competitors sales), and newsletter (number of newsletter subscription). There are various statistical methods for feature selection in a regression task, such as Chi-square Test, Variance Threshold, Dispersion Ratio, Random Forest Importance and etc. I used Correlation Coefficient to select predictors. Figure 1 presents the correlation matrix. Here, the variables with Coefficient less than 0.4 were dropped from the dataframe. In the final dataset only media1_S, competitor_sales, and newsletter are used to predict revenue.

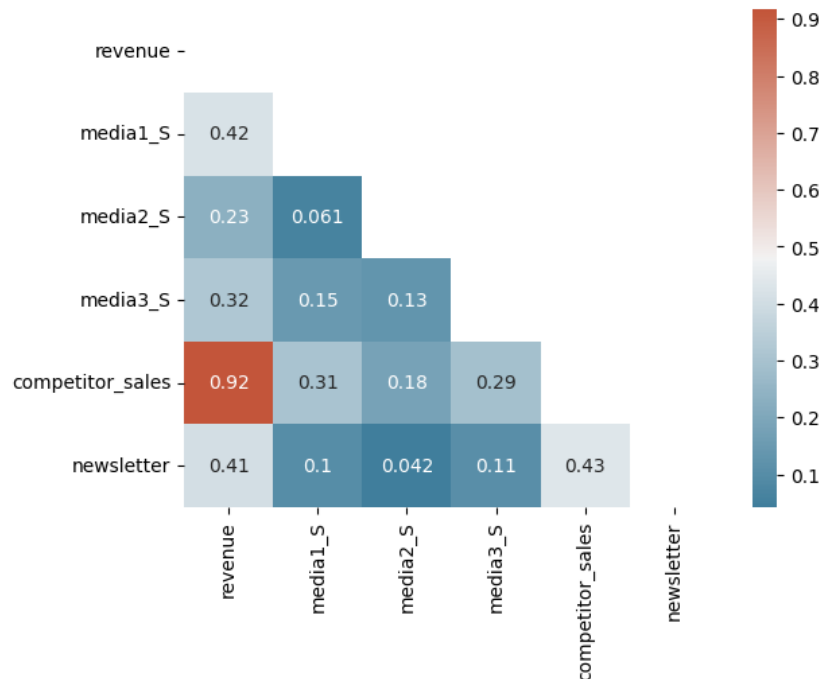


Figure 1: Correlation Matrix

For selection of the model, several commonly used regression model are used to model the marketing dataset. Using SKlearn library in Python, five algorithm are used for modelling the marketing data: linear regression, stochastic gradient descent regression, Random Forest, gradient boosting regression, Multilayer Perceptron (MLP). The results are presented in Table 1 and Figure 2.

Linear Regression and Stochastic Gradient Descent Regression models perform rather well in this task. The reported R^2 value for both models are 95%. Other methods does not perform well, but it is worth noting that each of these algorithms need hyperparameter tuning to perform well. Here, grid search was performed, but to reach the maximum potential of each model more analysis needs to be done. Analysing the effect of variables on

Table 1: Forecasting Error for Methods Used

Model	RMSE	MAE	R2	MAPE
Linear Regression	41441.49	34845.88	0.952	1.36
Stochastic Gradient Descent	41650.05	34783.47	0.952	1.36
Multi-layer Perceptron	123783.20	90774.79	0.576	3.55
Random Forest	139584.86	116637.73	0.461	4.79
Gradient Boosting	145422.27	136118.76	0.415	5.47

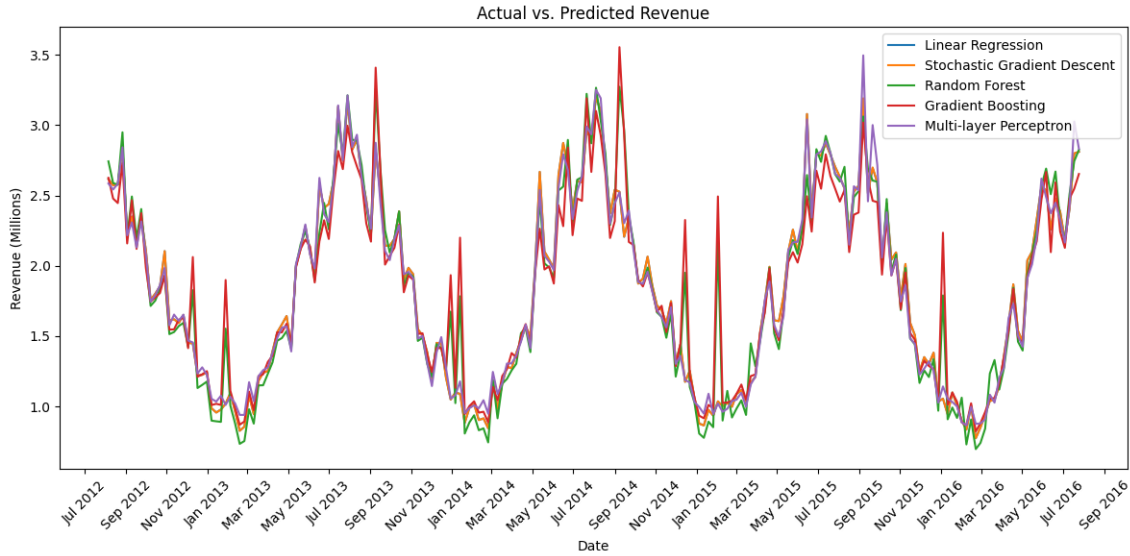


Figure 2: Forecasted Revenue

the regression results is an essential step in understanding the relationship between the independent variables (features) and the dependent variable (target) in a regression model. Some common methods are coefficient estimates, hypothesis tests, and residual analysis which will be performed in the next section. These are the coefficient of the regression model: 112513.37, 611400.75, and 12931.45.

The biggest limitation of the model used in the context of media marketing is the lack of model interpretability. Advanced regression models, like neural networks, can be complex and lack interpretability, making it challenging for marketers to understand the model's decision-making process. That's why the methods such as Bayesian model estimation can be useful in this field.

1.2 Prediction

This section overlaps with the previous section, the data was split into two datasets and appropriate variables were used for developing the prediction model. Figure 3 presents the forecasted revenue for the test period.

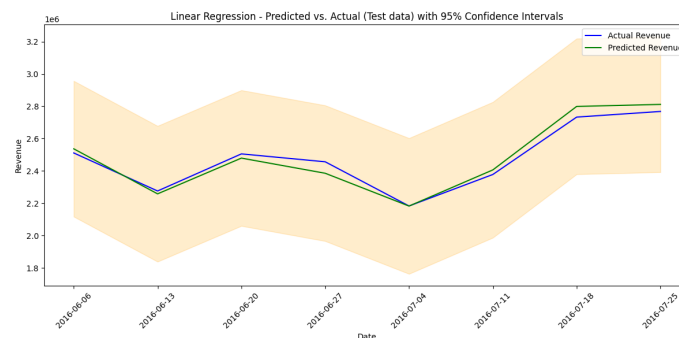


Figure 3: Forecast Results for Test Set

The performance of the model was evaluated using various metrics, RMSE, MAE, R^2 and MAPE. The results for the test set are presented in Table 1.

1.3 Temporal effects

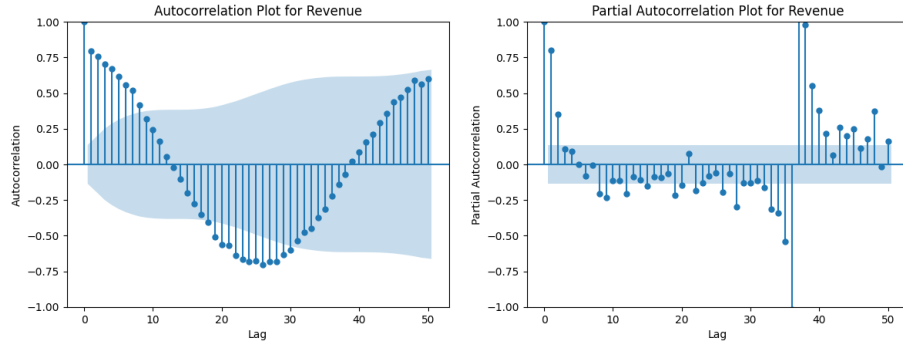


Figure 4: ACF and PACF Plots

As it is suggested by a colleague in the task, there are a clear affects of time seen in the marketing dataset. This can be seen by plotting the revenue against time. However, to be more precise, we can plot the Autocorrelation Function (ACF) of the revenue to determine whether trends and seasonal patterns are present. Figure 4 presents the ACF plot. As the autocorrelation is significant for number of lags, we can conclude that there is a temporal effect in the dataset. From the plot is can also be concluded that there is a clear trend and seasonality in the dataset. We can investigate this further by decomposition of the dataset. Figure 5 shows the result of the time series decomposition of the revenue. This enables us to identify patterns and trends that might impact the revenue data.

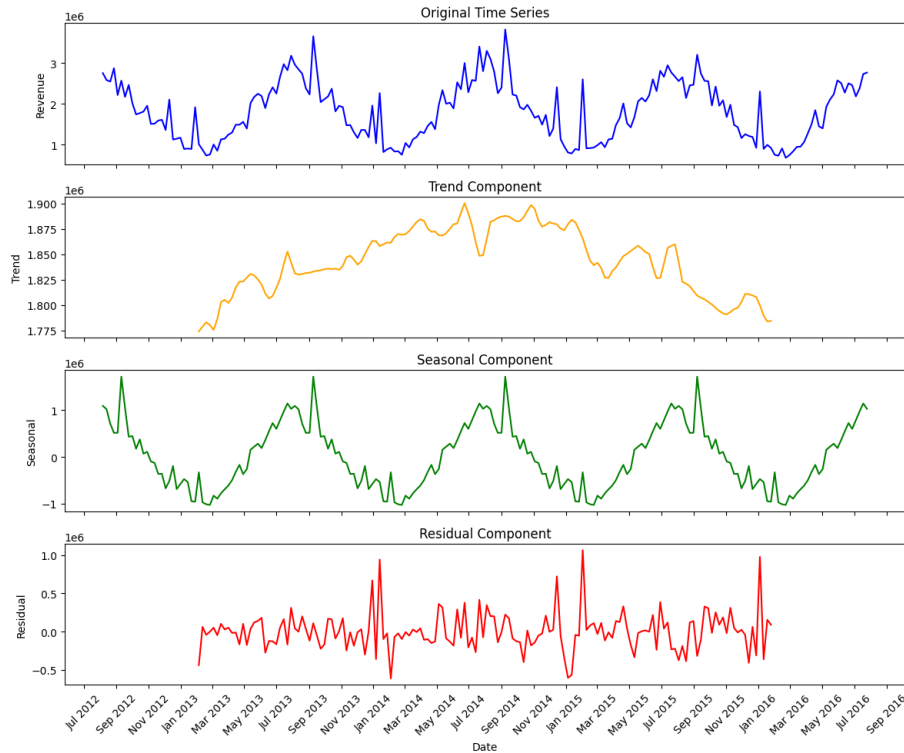


Figure 5: Time Series Decomposition of the Revenue

Next, we are going to model the dataset with the temporal components of the data. For this I used ARIMA to forecast the revenue. For this, we need to first check to see if the data is stationary, as this is one the requirements of the ARIMA algorithm. If the data is not stationary, we need to perform transformation such as differencing on the revenue. For selecting the autoregressive term or the number of lag observations in

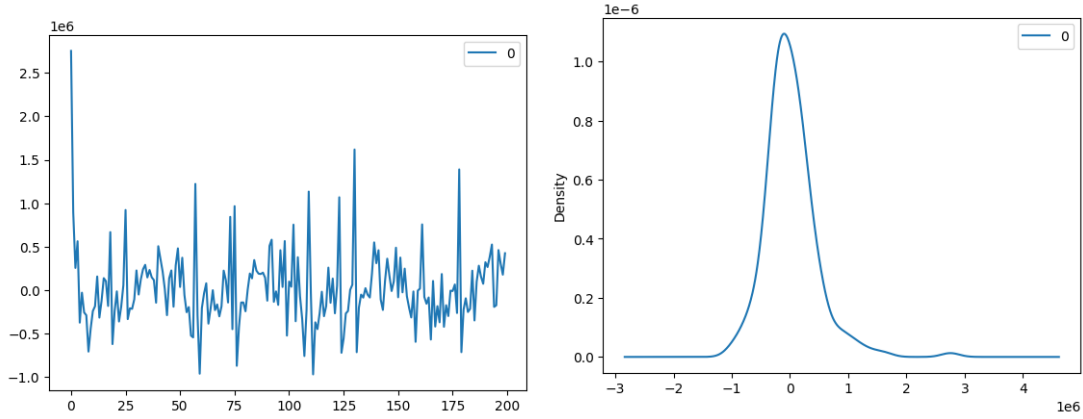


Figure 6: Line Plot and Denesity Plot of Residual Error

the ARIMA we need to plot the Partial Autocorrelation function (PACF) which is presented in Figure 4. By analysing the ACF and PACF the number of lag was selected as 12. This value was also selected during the Grid Search process for hyperparameter optimisation. Details of this are presented in the code. The best ARIMA order (p, d, q) is selected as following: 12, 1, 2. For this the RMSE value was reported as 217787.24.

Aside from the tipycal metrics that are used for evaluating the regression error, we can also plot the residuals as a line graph and also plot the density of residuals to check if the model has captured all of patterns in the data. This is presented in Figure 6.

Using ARIMA, we considered the temporal aspect of the data, but we discard the other features in the data. We can use other algorithms to predict revenue using all of the available features along side the temporal features. This can be done by using feature engineering, we can add new features to the data, like number of week, month, and year to do this. However, using windowing method we can transform the data and keep the temporal aspect of the data too. I wrote a function in python to do this. After this we can use any algorithm to perform prediction.

In the code submitted with this report a deep learning model is used to predict the revenue. The RMSE error is 213412.58, but please bear in mind that deep learning needs a large quantity of data for the algorithm to perform. Here we only used 208 data points. Also, by spending more time on the model and its architecture we can improve the result further. Figure 7 shows the result of the modelling for test set.

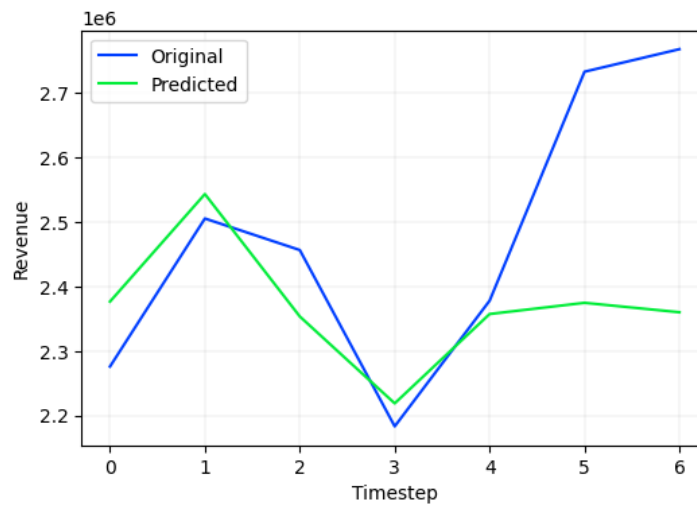


Figure 7: Deep Learning Results

2 Bayesian Methods

In this section I fit a simple neural network in a Bayesian framework using non-informative priors. In this example, the priors will be non-informative, which means they will have minimal impact on the final results and will let the data speak for itself. Non-informative priors are often used when we have little or no prior knowledge about the parameters we're trying to estimate, which is a good starting point in our Bayesian analysis.

Common noninformative priors include $unif(-1000, 1000)$ and $N(0, 10, 000)$. In this task I used a normal distribution for the priors, specifically, for weight and biases of the model, I used a normal distribution with mean zero and standard deviation of 10.

A prior distribution is said to be conjugate to a likelihood function if the resulting posterior distribution is in the same family of probability distributions as the prior. For example, in the case of a normal likelihood function with unknown mean and known variance, the conjugate prior is another normal distribution. If the likelihood is a binomial distribution, the conjugate prior is a beta distribution. If the likelihood is a Poisson distribution, the conjugate prior is a gamma distribution. Looking at the posterior distribution I can say the priors are conjugate. For MCMC I used the default value in PyMC, which is NUTS, a sampler for continuous variables based on Hamiltonian mechanics. It allows us to sample from the posterior distribution and obtain estimates for the model parameters.

Comparing the results from the first section and the result of the Bayesian framework produced similar outputs. The model has the RMSE value of 94420.70 which is on par with the performance of the previous model used. Figure ?? shows the result of the prediction on the test set.

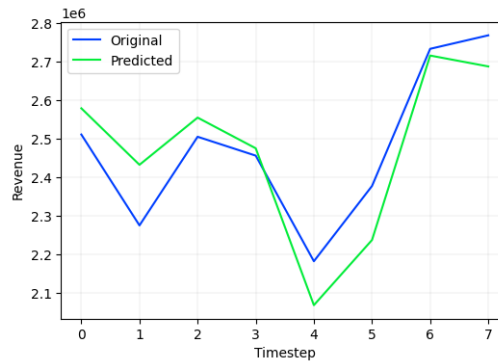


Figure 8: Bayesian Neural Network Result for Test Set

2.1 Prior choice

To modify the prior specification to account for the expert's belief that "Media 3 has no impact in generating marketing revenues," we can update the prior for the corresponding coefficient to reflect this belief. We want to incorporate the expert's opinion while still allowing the data to influence the final estimates. We can do this by using a prior distribution for the coefficient of "Media 3" with a high variance, indicating a wide range of possible values. For example, we can use a Normal distribution for the coefficient of Media 3 with a mean of 0 and a large standard deviation, such as 10. This would express the expert's belief that the coefficient is likely to be around 0, but they are not certain and allow the data to adjust this belief.

However, I already use a normal distribution with mean 0 and standard deviation of 10 to accommodate non informative priors in the model. So to answer the question raised in the task, I must say based on the estimated parameters from the model, the expert is correct and Media 3 has no to little impact on the revenue.

Additionally, to incorporate the opinions of the other expert, we can set informative priors for the coefficients of "newsletter," "media1_S," "media2_S," and "competitor_sales" based on the expert's beliefs. For example:

"newsletter": We can use a Normal distribution with a mean of 1 and a small standard deviation, such as 0.1, to express the belief that there is a strictly positive relationship between "newsletter" and marketing revenues. For "media1_S" we can use a Normal distribution with a mean of 2 and a small standard deviation, such as 0.1, to express the belief that one unit invested in "Media 1" will result in two units of marketing revenues returned. For "media2_S" we can use a Normal distribution with a mean of 8 (four times the mean for "media1_S") and a small standard deviation, such as 0.1, to express the belief that the impact of "Media

2" is four times that of "Media 1."

For "competitor_sales" we can use a Normal distribution with a mean of 0.15 and a small standard deviation, such as 0.05, to express the belief that for every unit increase in "competitor_sales," the change on revenues returned must range between 0 and 0.3.

Each of these hypothesis need to be implemented in the model separately and the results and posterior distribution needs to be carefully study. Unfortunately, I did not had enough time to do this analysis. The code needed for doing this is provided in the submitted code. Given more time, I could have make the necessary analysis to come up with answers.

3 Closing Remarks

I would want to convey my gratitude for the opportunity to demonstrate my qualifications for this post at the University of Essex and Croud Inc. This is an intriguing role for me, and I would be pleased to go to the next stage of the hiring procedure.

The Python codes can be found on the following pages. You will also receive a link to a colab notebook along with this report.

[Click here to open in Google Colab](#)

Python Code

```
1  # -*- coding: utf-8 -*-
2  """alotfipoor_ktp_associate_essex.ipynb
3
4  Automatically generated by Colaboratory.
5
6  Original file is located at
7      https://colab.research.google.com/drive/16cX08SFB-SyHqLJpUx33vgg6iqRHmVel
8
9  # Setup
10 """
11
12 # Commented out IPython magic to ensure Python compatibility.
13 # from datetime import datetime
14 import os, glob, re
15 import numpy as np
16 import pandas as pd
17 import warnings
18 from math import sqrt
19
20 warnings.simplefilter(action='ignore', category=FutureWarning)
21 warnings.simplefilter(action='ignore', category=DeprecationWarning)
22
23 # import seaborn as sns
24 import matplotlib as mpl
25 import matplotlib.pyplot as plt
26 # %matplotlib inline
27 import seaborn as sns
28 from matplotlib.ticker import FuncFormatter
29 import matplotlib.dates as mdates
30
31 from IPython.core.display import display, HTML
32 display(HTML("<style>.container { width:100% !important; }</style>"))
33
34 # plotly
35 import plotly.express as px
36 import plotly.graph_objs as go
37 import plotly.io as pio
38 pio.renderers.default = 'colab'
39
40 from statsmodels.tsa.stattools import adfuller
41 from statsmodels.tsa.arima.model import ARIMA
42 from statsmodels.graphics.tsaplots import plot_pacf
43 from statsmodels.graphics.tsaplots import plot_acf
44 from statsmodels.tsa.seasonal import seasonal_decompose
45
46 import pymc as pm
47 import arviz as az
48
49 # import optuna
50
51 # from scipy.stats import uniform, randint
52 # from sklearn.model_selection import TimeSeriesSplit, cross_val_score, GridSearchCV, RandomizedSearchCV
53
54 from sklearn.model_selection import train_test_split
55 from sklearn.preprocessing import StandardScaler, MinMaxScaler
56 from sklearn.linear_model import LinearRegression, SGDRegressor
57 from sklearn.svm import SVR
58 from sklearn.gaussian_process import GaussianProcessRegressor
59 from sklearn.gaussian_process.kernels import WhiteKernel, DotProduct
60 from sklearn.tree import DecisionTreeRegressor
```



```
61 from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
62 from sklearn.neural_network import MLPRegressor
63 from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score, mean_squared_log_error
64
65 from keras.models import Sequential
66 from keras.layers import Dense, LSTM, Dropout, TimeDistributed, Flatten, BatchNormalization
67 from tensorflow.keras.optimizers import Adam
68 from tensorflow.keras.callbacks import LearningRateScheduler, EarlyStopping, ReduceLROnPlateau, ModelCheckpoint
69 # import tensorflow as tf
70
71 # data Table Display
72 # %load_ext google.colab.data_table
73
74 # to make this notebook's output stable across runs
75 np.random.seed(42)
76
77 # mount google colab
78 from google.colab import drive
79 drive.mount('/content/drive')
80
81 # load dataset
82 df = pd.read_csv('/content/drive/MyDrive/[04] Colab Notebooks/data/weekly_media_sample.csv')
83 df['DATE'] = pd.to_datetime(df['DATE'])
84
85 df.head(5)
86
87 fig = go.Figure()
88 fig.add_trace(go.Scatter(x=df.DATE, y=df.revenue, name='Revenue', line=dict(color='blue', width=1)))
89 fig.add_trace(go.Scatter(x=df.DATE, y=df.competitor_sales, name='Competitor Sales', line=dict(color='red',
90 ↪ width=.8)))
91 fig.add_trace(go.Scatter(x=df.DATE, y=df.newsletter, name='Newsletter Subscription', line=dict(color='green',
92 ↪ width=.8)))
93 fig.show()
94
95 """# Part 1: Classical Marketing Data Modelling"""
96
97 # Creating a subset of dataframe for modelling
98 # df2 = df[['DATE', 'revenue', 'media1_S', 'media2_S', 'media3_S', 'competitor_sales', 'newsletter']].copy()
99 df2 = df[['DATE', 'revenue', 'media1_S', 'competitor_sales', 'newsletter']].copy()
100
101 # Calculate correlation coefficients with the target variable 'revenue'
102 correlation_matrix = df2.corr()
103 correlation_with_revenue = correlation_matrix['revenue'].abs().sort_values(ascending=False)
104
105 # Print the correlation coefficients
106 print(correlation_with_revenue)
107
108 mask = np.triu(np.ones_like(correlation_matrix, dtype=bool))
109 cmap = sns.diverging_palette(230, 20, as_cmap=True)
110 sns.heatmap(correlation_matrix, mask=mask, cmap=cmap, annot=True)
111
112 # Define the split date based on the given threshold (e.g., '2016-05-30' in this case)
113 split_date = '2016-05-30'
114
115 # Split the data into train and test sets
116 train = df2[df2.DATE <= split_date]
117 test = df2[df2.DATE > split_date]
118
119 # Extract the target variable 'revenue' from the training and test data
120 y_train = train['revenue']
121 y_test = test['revenue']
122
123 # Extract the input features from the training and test data
```

```
122 X_train = train.drop(['revenue', 'DATE'], axis=1)
123 X_test = test.drop(['revenue', 'DATE'], axis=1)
124
125 # Scale the data
126 scaler = StandardScaler()
127 X_train = scaler.fit_transform(X_train)
128 X_test = scaler.transform(X_test)
129
130 # Create instances of regression models with hyperparameters
131 linear_reg = LinearRegression()
132 sgdr_reg = SGDRegressor()
133 # svm_reg = SVR(C=10, degree=3, gamma='auto', kernel='rbf')
134 # gaussian_process_reg = GaussianProcessRegressor(kernel=DotProduct() + WhiteKernel())
135 random_forest_reg = RandomForestRegressor(n_estimators=200, max_depth=20, random_state=42)
136 gradient_boosting_reg = GradientBoostingRegressor(n_estimators=200, learning_rate=0.01, max_depth=20,
137 ↪ random_state=42)
138
139 mlp_reg = MLPRegressor(hidden_layer_sizes=(500, 300, 200), batch_size=64, learning_rate='adaptive', alpha=0.1,
140 ↪ max_iter=5000, random_state=42)
141
142 # Create DataFrames to store predicted values
143 train_predictions_df = pd.DataFrame(data=y_train.values, columns=['Actual Revenue'])
144 test_predictions_df = pd.DataFrame(data=y_test.values, columns=['Actual Revenue'])
145
146 results_df = pd.DataFrame(columns=['Model', 'RMSE', 'MAE', 'R2', 'MAPE'])
147
148 # Mean Absolute Percentage Error (MAPE) function
149 def mean_absolute_percentage_error(y_true, y_pred):
150     return np.mean(np.abs((y_true - y_pred) / y_true)) * 100
151
152 # Function to evaluate and append results to the DataFrame
153 def evaluate_model(model_name, model, X_train, y_train, X_test, y_test):
154     model.fit(X_train, y_train)
155     y_train_pred = model.predict(X_train)
156     y_test_pred = model.predict(X_test)
157     train_predictions_df[model_name] = y_train_pred
158     test_predictions_df[model_name] = y_test_pred
159     rmse = mean_squared_error(y_test, y_test_pred, squared=False)
160     mae = mean_absolute_error(y_test, y_test_pred)
161     r2 = r2_score(y_test, y_test_pred)
162     mape = mean_absolute_percentage_error(y_test, y_test_pred)
163     results_df.loc[len(results_df)] = [model_name, rmse, mae, r2, mape]
164
165 # Evaluate each model and store the results
166 evaluate_model('Linear Regression', linear_reg, X_train, y_train, X_test, y_test)
167 evaluate_model('Stochastic Gradient Descent', sgdr_reg, X_train, y_train, X_test, y_test)
168 # evaluate_model('SVM', svm_reg, X_train, y_train, X_test, y_test)
169 # evaluate_model('Gaussian Processes', gaussian_process_reg, X_train, y_train, X_test, y_test)
170 evaluate_model('Random Forest', random_forest_reg, X_train, y_train, X_test, y_test)
171 evaluate_model('Gradient Boosting', gradient_boosting_reg, X_train, y_train, X_test, y_test)
172 evaluate_model('Multi-layer Perceptron', mlp_reg, X_train, y_train, X_test, y_test)
173
174 results_df.head()
175
176 # Print the coefficient of the linear regression
177 linear_reg.coef_
178
179 # Merge train_predictions_df and test_predictions_df while keeping 'DATE' column
180 merged_predictions_df = pd.concat([train_predictions_df, test_predictions_df])
181
182 # Add the 'DATE' column from df2 to merged_predictions_df
183 merged_predictions_df['DATE'] = np.sort(df2['DATE'])
184
185 # Create a trace for each model's predicted values against the actual 'revenue'
```

```
183 traces = []
184 for column in merged_predictions_df.columns[:-1]: # Exclude the first and last columns ('Actual Revenue' and
↪ 'DATE')
185     trace = go.Scatter(
186         x=merged_predictions_df['DATE'],
187         y=merged_predictions_df[column],
188         mode='lines',
189         name=column
190     )
191     traces.append(trace)
192
193 # Create the layout for the plot
194 layout = go.Layout(
195     title='Actual vs. Predicted Revenue',
196     xaxis=dict(title='Date'),
197     yaxis=dict(title='Revenue'),
198 )
199
200 # Create the figure and plot
201 fig = go.Figure(data=traces, layout=layout)
202 fig.show()
203
204 # Convert 'DATE' column to datetime object
205 merged_predictions_df['DATE'] = pd.to_datetime(merged_predictions_df['DATE'])
206
207 # Create a line chart for each model's predicted values against the actual 'revenue' over time
208 plt.figure(figsize=(12, 6))
209 for column in merged_predictions_df.columns[1:-1]: # Exclude the first and last columns ('Actual Revenue' and
↪ 'DATE')
210     plt.plot(merged_predictions_df['DATE'], merged_predictions_df[column], label=column)
211
212 # Set the x-axis label to 'Date' and y-axis label to 'Revenue'
213 plt.xlabel('Date')
214 plt.ylabel('Revenue (Millions)')
215
216 # Set the title of the plot
217 plt.title('Actual vs. Predicted Revenue')
218
219 # Show the legend to identify each model's line
220 plt.legend()
221
222 # Format x-axis date labels to show only the month and year
223 date_format = mdates.DateFormatter("%b %Y")
224 plt.gca().xaxis.set_major_formatter(date_format)
225
226 # Set the x-axis tick interval to show labels every n months (adjust n as needed)
227 n_months = 2
228 plt.gca().xaxis.set_major_locator(mdates.MonthLocator(interval=n_months))
229
230 # Rotate the x-axis labels for better visibility
231 plt.xticks(rotation=45)
232
233 # Format y-axis tick labels in millions
234 def millions_formatter(x, pos):
235     return f'{x / 1e6:.1f}'
236
237 formatter = FuncFormatter(millions_formatter)
238 plt.gca().yaxis.set_major_formatter(formatter)
239
240 # Show the plot
241 plt.tight_layout()
242 plt.show()
243
```

```
244 # Fit the Linear Regression model
245 linear_reg.fit(X_train, y_train)
246
247 # Predict on the test set
248 Y_pred_linear = linear_reg.predict(X_test)
249
250 # Calculate the 95% confidence intervals for the predictions
251 conf_interval_linear = 1.96 * np.std(Y_pred_linear) # 1.96 is the z-score for 95% confidence
252
253 # # Create a DataFrame with predictions and confidence intervals for the last 20 rows
254 # df_pred_linear = pd.DataFrame({'DATE': df2['DATE'].tail(8).values,
255 #                               'Actual Revenue': y_test[-8:].values.flatten(),
256 #                               'Predicted Revenue': Y_pred_linear[-8:].flatten(),
257 #                               'Lower CI': (Y_pred_linear[-8:] - conf_interval_linear).flatten(),
258 #                               'Upper CI': (Y_pred_linear[-8:] + conf_interval_linear).flatten()})
259
260 # Create a DataFrame with predictions and confidence intervals for the last 20 rows
261 df_pred_linear = pd.DataFrame({'DATE': df2['DATE'].tail(8).values,
262                               'Actual Revenue': y_test.values.flatten(),
263                               'Predicted Revenue': Y_pred_linear.flatten(),
264                               'Lower CI': (Y_pred_linear - conf_interval_linear).flatten(),
265                               'Upper CI': (Y_pred_linear + conf_interval_linear).flatten()})
266
267 # Plot the last 20 rows of predictions, actual observations, and confidence intervals
268 plt.figure(figsize=(12, 6))
269 plt.plot(df2['DATE'].tail(8), df2['revenue'].tail(8), label='Actual Revenue', color='blue')
270 plt.plot(df_pred_linear['DATE'], df_pred_linear['Predicted Revenue'], label='Predicted Revenue', color='green')
271 plt.fill_between(df_pred_linear['DATE'], df_pred_linear['Lower CI'], df_pred_linear['Upper CI'], alpha=0.2,
272                ↪ color='orange')
273 plt.xlabel("Date")
274 plt.ylabel("Revenue")
275 plt.title("Linear Regression - Predicted vs. Actual (Test data) with 95% Confidence Intervals")
276 plt.legend()
277 plt.xticks(rotation=45)
278 plt.tight_layout()
279 plt.show()
280
281 """"# ARIMA""""
282
283 sns_c = sns.color_palette(palette="deep")
284
285 sns.pairplot(
286     data=df2, kind="scatter", height=2, plot_kws={"color": sns_c[1]}, diag_kws={"color": sns_c[2]}
287 );
288
289 plt.figure(figsize=(12, 6))
290 plot_acf(df2['revenue'], lags=50)
291 plt.xlabel('Lag')
292 plt.ylabel('Autocorrelation')
293 plt.title('Autocorrelation Plot for Revenue')
294 plt.show()
295
296 plt.figure(figsize=(12, 6))
297 plot_pacf(df2['revenue'], lags=50)
298 plt.xlabel('Lag')
299 plt.ylabel('Partial Autocorrelation')
300 plt.title('Partial Autocorrelation Plot for Revenue')
301 plt.show()
302
303 # Performing seasonal decomposition
304 result = seasonal_decompose(df2['revenue'], model='additive', period=52) # Assuming weekly seasonality
305 ↪ (period=52)
```

```
305 # Creating a subplot layout
306 fig, axes = plt.subplots(4, 1, figsize=(12, 10), sharex=True)
307
308 # Plotting the original time series
309 axes[0].plot(df2['DATE'], df2['revenue'], label='Original', color='blue')
310 axes[0].set_ylabel('Revenue')
311 axes[0].set_title('Original Time Series')
312
313 # Plotting the trend component
314 axes[1].plot(df2['DATE'], result.trend, label='Trend', color='orange')
315 axes[1].set_ylabel('Trend')
316 axes[1].set_title('Trend Component')
317
318 # Plotting the seasonal component
319 axes[2].plot(df2['DATE'], result.seasonal, label='Seasonal', color='green')
320 axes[2].set_ylabel('Seasonal')
321 axes[2].set_title('Seasonal Component')
322
323 # Plotting the residual component
324 axes[3].plot(df2['DATE'], result.resid, label='Residual', color='red')
325 axes[3].set_xlabel('Date')
326 axes[3].set_ylabel('Residual')
327 axes[3].set_title('Residual Component')
328
329 # Format x-axis date labels to show only the month and year
330 date_format = mdates.DateFormatter("%b %Y")
331 axes[-1].xaxis.set_major_formatter(date_format)
332
333 # Set the x-axis tick interval to show labels every n months (adjust n as needed)
334 n_months = 2
335 axes[-1].xaxis.set_major_locator(mdates.MonthLocator(interval=n_months))
336
337 # Rotate the x-axis labels for better visibility
338 plt.xticks(rotation=45)
339
340 # Adjusting the layout
341 plt.tight_layout()
342
343 # Show the plot
344 plt.show()
345
346 # Perform the ADF test on the 'revenue' data
347 result = adfuller(df2['revenue'])
348
349 # Extract the test statistic and p-value
350 test_statistic = result[0]
351 p_value = result[1]
352
353 print("ADF Test Statistic:", test_statistic)
354 print("p-value:", p_value)
355
356 if p_value <= 0.05:
357     print("The data is stationary.")
358 else:
359     print("The data is not stationary.")
360
361 # Find the optimal ARIMA hyperparameters using grid search
362 best_aic = float("inf")
363 best_order = None
364
365 # Define the ranges for p, d, and q
366 p_range = range(0, 15) # Example: p can be 0, 1, or 2
367 d_range = range(0, 2) # Example: d can be 0 or 1
```

```
368 q_range = range(0, 3) # Example: q can be 0, 1, or 2
369
370 for p in p_range:
371     for d in d_range:
372         for q in q_range:
373             try:
374                 arima_model = ARIMA(train['revenue'], order=(p, d, q))
375                 arima_fit = arima_model.fit()
376
377                 # Calculate AIC score for the current ARIMA model
378                 aic = arima_fit.aic
379
380                 if aic < best_aic:
381                     best_aic = aic
382                     best_order = (p, d, q)
383
384             except:
385                 continue
386
387 # Fit the ARIMA model with the best parameters to the training data
388 arima_model = ARIMA(train['revenue'], order=best_order)
389 arima_fit = arima_model.fit()
390
391 # Make predictions on the test set
392 y_pred_arima = arima_fit.forecast(steps=len(test))
393
394 # Evaluate the performance using Mean Squared Error (MSE)
395 mse_arima = mean_squared_error(test['revenue'], y_pred_arima)
396 rmse_arima = np.sqrt(mse_arima)
397
398 print("Best ARIMA Order (p, d, q):", best_order)
399 print("ARIMA RMSE:", rmse_arima)
400
401 # Plot the actual revenue and ARIMA predictions
402 plt.figure(figsize=(10, 6))
403 plt.plot(df2.index, df2['revenue'], label='Actual Revenue', color='blue')
404 plt.plot(test.index, y_pred_arima, label='ARIMA Predictions', color='red')
405 plt.xlabel('Date')
406 plt.ylabel('Revenue')
407 plt.title('Actual Revenue vs. ARIMA Predictions')
408 plt.legend()
409 plt.show()
410
411 # summary of fit model
412 print(arima_fit.summary())
413
414 # line plot of residuals
415 residuals = pd.DataFrame(arima_fit.resid)
416 residuals.plot()
417 plt.show()
418
419 # density plot of residuals
420 residuals.plot(kind='kde')
421 plt.show()
422 # summary stats of residuals
423 print(residuals.describe())
424
425 """## Deep learning"""
426
427 def create_time_series_dataset_with_lags(df, threshold_date, lag_observation=5, forecast_horizon=1):
428     """
429     Create X and Y matrices for time series modeling with lag observations and split into train and test sets.
430
```

```
431 Parameters:
432 df (pd.DataFrame): DataFrame containing the time series data with columns: revenue, competitor_sales,
433 newsletter, Total_Media_Spend, and DATE.
434 threshold_date (str): The date to split the data into train and test sets (format: 'YYYY-MM-DD').
435 lag_observation (int): The number of lag observations to use. Default is 5.
436 forecast_horizon (int): The number of steps ahead to forecast revenue. Default is 1.
437
438 Returns:
439 X_train (np.ndarray): 3-dimensional array containing input features with lag observations for the train
↪ set.
440 Y_train (np.ndarray): 1-dimensional array containing target values (revenue) for the train set.
441 X_test (np.ndarray): 3-dimensional array containing input features with lag observations for the test
↪ set.
442 y_test (np.ndarray): 1-dimensional array containing target values (revenue) for the test set.
443 """
444 # Convert threshold_date to Timestamp object
445 threshold_date = pd.to_datetime(threshold_date)
446
447
448 # Get the 'revenue' time series and other features
449 revenue = df['revenue'].values
450 competitor_sales = df['competitor_sales'].values
451 newsletter = df['newsletter'].values
452 media1_S = df['media1_S'].values
453
454 X_train, Y_train, X_test, Y_test = [], [], [], []
455
456 for i in range(len(df) - lag_observation - forecast_horizon + 1):
457     # Input sequence with lag observations for all features
458     x_sequence = np.column_stack((
459         revenue[i:i + lag_observation],
460         competitor_sales[i:i + lag_observation],
461         newsletter[i:i + lag_observation],
462         media1_S[i:i + lag_observation]
463     ))
464
465     if df['DATE'][i + lag_observation - 1] <= threshold_date:
466         # Add to train set
467         X_train.append(x_sequence)
468         Y_train.append(revenue[i + lag_observation + forecast_horizon - 1])
469     else:
470         # Add to test set
471         X_test.append(x_sequence)
472         Y_test.append(revenue[i + lag_observation + forecast_horizon - 1])
473
474 X_train = np.array(X_train)
475 Y_train = np.array(Y_train)
476 X_test = np.array(X_test)
477 Y_test = np.array(Y_test)
478
479 return X_train, Y_train, X_test, Y_test
480
481 threshold_date = '2016-05-30'
482 lag_observation = 3
483 forecast_horizon = 1
484
485 X_train, Y_train, X_test, Y_test = create_time_series_dataset_with_lags(df2, threshold_date, lag_observation,
↪ forecast_horizon)
486
487 # Normalize the input data
488 scaler = StandardScaler()
489 X_train_scaled = scaler.fit_transform(X_train.reshape(-1, X_train.shape[-1])).reshape(X_train.shape)
490 X_test_scaled = scaler.transform(X_test.reshape(-1, X_test.shape[-1])).reshape(X_test.shape)
```

```
491
492 # Define the model
493 model = Sequential()
494 model.add(LSTM(128, activation='relu', return_sequences=True, input_shape=(X_train.shape[1], X_train.shape[2])))
495 model.add(LSTM(64, activation='relu', return_sequences=True))
496 model.add(Dropout(0.2))
497 model.add(Flatten())
498 model.add(Dense(128))
499 model.add(Dense(64))
500 model.add(Dense(64))
501 model.add(Dense(1))
502
503 model.compile(optimizer='adam', loss='mean_squared_error')
504
505 # Train the model
506 history = model.fit(X_train_scaled, Y_train, epochs=200, batch_size=16, verbose=0, shuffle=False)
507
508 # Make predictions
509 Y_pred = model.predict(X_test_scaled)
510
511 rmse = mean_squared_error(Y_test, Y_pred, squared = False)
512 print('RMSE for test dataset = %.5f' % rmse)
513
514 df_predicted = pd.DataFrame({'Actual Revenue': Y_test.flatten(), 'Predicted Revenue': Y_pred.flatten()})
515 print(df_predicted)
516
517 with plt.style.context('seaborn-bright', after_reset=True):
518     plt.figure(figsize=(6,4))
519     plt.plot(Y_test, label='Original')
520     plt.plot(Y_pred, label='Predicted')
521     plt.legend()
522     plt.xlabel('Timestep')
523     plt.ylabel('Revenue')
524     plt.grid(color='gray', linestyle='-', linewidth=0.1)
525     plt.show()
526
527 """# Part 2: Bayesian methods"""
528
529 # Define the split date based on the given threshold (e.g., '2016-05-30' in this case)
530 split_date = '2016-05-30'
531
532 # Split the data into train and test sets
533 train = df[df.DATE <= split_date]
534 test = df[df.DATE > split_date]
535
536 # Extract the target variable 'revenue' from the training and test data
537 y_train = train['revenue']
538 y_test = test['revenue']
539
540 # Extract the input features from the training and test data
541 x_train = train.drop(['revenue', 'DATE', 'X'], axis=1)
542 x_test = test.drop(['revenue', 'DATE', 'X'], axis=1)
543
544 # Scale the data
545 scaler = StandardScaler()
546 X_train = scaler.fit_transform(x_train)
547 X_test = scaler.transform(x_test)
548
549 # Define different priors for each feature's weights
550 priors = {
551     'media1_S': {
552         'mu': 0,
553         'sd': 10
```



```

554     },
555     'media2_S': {
556         'mu': 0,
557         'sd': 10
558     },
559     'media3_S': {
560         'mu': 0,
561         'sd': 10
562     },
563     'competitor_sales': {
564         'mu': 0,
565         'sd': 10
566     },
567     'newsletter': {
568         'mu': 0,
569         'sd': 10
570     }
571 }
572
573 def construct_bnn(x_input, y_input):
574     # Using the context manager to build the model
575     with pm.Model() as bnn:
576         # Data input
577         x_data = pm.MutableData('x_data', x_input)
578         y_data = pm.MutableData('y_data', y_input)
579
580         # Model structure
581         num_hidden = 5
582
583         # Priors for the weights
584         weights = {}
585
586         for feature in x_train.columns:
587             weights[feature] = pm.Normal(f'w_{feature}', mu=priors[feature]['mu'], sigma=priors[feature]['sd'],
588                                         ↪ shape=num_hidden)
589
590         b_1 = pm.Normal('b_1', mu=0, sigma=10, shape=num_hidden)
591         w_out = pm.Normal('w_out', mu=0, sigma=10, shape=num_hidden)
592         b_out = pm.Normal('b_out', mu=0, sigma=10, shape=1)
593
594         # Neural network architecture
595         act_1 = pm.math.tanh(pm.math.dot(x_data, pm.math.stack([weights[feature] for feature in
596                                         ↪ x_train.columns]))) + b_1)
597
598         act_out = pm.Deterministic('act_out', pm.math.dot(act_1, w_out) + b_out)
599
600         # Likelihood of the data (observed)
601         sigma = pm.HalfNormal('sigma', sigma=1)
602         likelihood = pm.Normal('Y_obs', mu=act_out, sigma=1, observed=y_data, total_size=y_input.shape[0])
603
604         return bnn
605
606 bnn = construct_bnn(X_train, y_train)
607 # pm.model_to_graphviz(bnn)
608
609 # Specifying the MCMC algorithm (NUTS)
610 with bnn:
611     step = pm.NUTS()
612     trace = pm.sample(1000, tune=1000, step=step, chains=4, random_seed=42)
613
614 pm.plot_trace(trace, figsize=(14,10), legend=True, compact=False)
615 plt.show()

```

```
615
616 pm.set_data({"x_data": X_test, "y_data": y_test}, model=bnn)
617
618 # Generate posterior samples.
619 ppc_test = pm.sample_posterior_predictive(trace, model=bnn)
620
621 # Compute the point prediction by taking the mean
622 y_test_pred = ppc_test.posterior_predictive['Y_obs'].mean(dim=['chain', 'draw'])
623
624 # y_mean = y_test_pred.mean(axis=0)
625 rmse = mean_squared_error(y_test, y_test_pred, squared = False)
626 print('RMSE for test dataset = %.5f' % rmse)
627
628 with plt.style.context('seaborn-bright', after_reset=True):
629     plt.figure(figsize=(6,4))
630     plt.plot(y_test.values, label='Original')
631     plt.plot(y_test_pred.values, label='Predicted')
632     plt.legend()
633     plt.xlabel('Timestep')
634     plt.ylabel('Revenue')
635     plt.grid(color='gray', linestyle='-', linewidth=0.1)
636     plt.show()
637
638 pm.plot_trace(trace, figsize=(14,10), legend=True, compact=False)
639 plt.show()
640
641 pm.summary(trace, round_to=2)
642
643 # Plotting the ACF plot for the MCMC chains
644 pm.plot_autocorr(trace)
645 plt.show()
646
647 az.plot_posterior(trace)
648
649 """## Prior Choice"""
650
651 # Define different priors for each feature's weights
652 priors = {
653     'media1_S': {
654         'mu': 2,
655         'sd': 0.1
656     },
657     'media2_S': {
658         'mu': 8,
659         'sd': 0.1
660     },
661     'media3_S': {
662         'mu': 0,
663         'sd': 10
664     },
665     'competitor_sales': {
666         'mu': 0.15,
667         'sd': 0.05
668     },
669     'newsletter': {
670         'mu': 1,
671         'sd': 0.1
672     }
673 }
674
675 def construct_bnn(x_input, y_input):
676     # Using the context manager to build the model
677     with pm.Model() as bnn:
```

```
678     # Data input
679     x_data = pm.MutableData('x_data', x_input)
680     y_data = pm.MutableData('y_data', y_input)
681
682     # Model structure
683     num_hidden = 1
684
685     # Priors for the weights
686     weights = {}
687
688     for feature in x_train.columns:
689         weights[feature] = pm.Normal(f'w_{feature}', mu=priors[feature]['mu'], sigma=priors[feature]['sd'],
690                                     ↪ shape=num_hidden)
691
692     b_1 = pm.Normal('b_1', mu=0, sigma=10, shape=num_hidden)
693     w_out = pm.Normal('w_out', mu=0, sigma=10, shape=num_hidden)
694     b_out = pm.Normal('b_out', mu=0, sigma=10, shape=1)
695
696     # Neural network architecture
697     act_1 = pm.math.tanh(pm.math.dot(x_data, pm.math.stack([weights[feature] for feature in
698     ↪ x_train.columns]))) + b_1)
699
700     act_out = pm.Deterministic('act_out', pm.math.dot(act_1, w_out) + b_out)
701
702     # Likelihood of the data (observed)
703     sigma = pm.HalfNormal('sigma', sigma=1)
704     likelihood = pm.Normal('Y_obs', mu=act_out, sigma=1, observed=y_data, total_size=y_input.shape[0])
705
706     return bnn
707
708 bnn = construct_bnn(X_train, y_train)
709 # pm.model_to_graphviz(bnn)
710
711 # Specifying the MCMC algorithm (NUTS)
712 with bnn:
713     step = pm.NUTS()
714     trace = pm.sample(1000, tune=1000, step=step, chains=4, random_seed=42)
715
716 pm.set_data({'x_data': X_test, "y_data": y_test}, model=bnn)
717
718 # Generate posterior samples.
719 ppc_test = pm.sample_posterior_predictive(trace, model=bnn)
720
721 # Compute the point prediction by taking the mean
722 y_test_pred = ppc_test.posterior_predictive['Y_obs'].mean(dim=['chain', 'draw'])
723
724 # y_mean = y_test_pred.mean(axis=0)
725 rmse = mean_squared_error(y_test, y_test_pred, squared = False)
726 print('RMSE for test dataset = %.5f' % rmse)
```